

# Snowflake SQL

Snowflake objects including users, virtual warehouses, databases, schemas, tables, views, columns, functions, and stored procedures are created, manipulated, and modified using **Data Definition Language (DDL)** commands. They're also used to do things like set parameters, initialize variables, and start transactions at the account and session levels.

All DDL commands are built on the foundation of the following commands:

- **ALTER**
- **COMMENT**
- **CREATE**
- **DESCRIBE**
- **DROP**
- **SHOW**
- **USE**

## Snowflake SQL: DML Commands

**DML (Data Manipulation Language)** commands are used to modify and update data. You can use DML commands to insert, delete, update, and merge data in Snowflake tables using the following commands:

- **INSERT**
- **INSERT (multi-table)**
- **MERGE**
- **UPDATE**
- **DELETE**
- **TRUNCATE TABLE**

## Snowflake SQL: Query Syntax

The most typical application of **SQL queries** is to discover particular data by filtering against predefined parameters. Data management operations can also be automated using them. For robust relational database querying, Snowflake's Data Cloud platform has a Data Warehouse workload that supports the most prevalent standardized version of SQL (ANSI). It can also

combine semi-structured data, such as JSON, with structured data stored in SQL. **Snowflake SQL** makes JSON data more accessible and enables users to mix it with structured data. Users may access JSON data using SQL queries and easily link it to traditional tabular data using Snowflake.

Querying in Snowflake can be done using conventional **SELECT** statements and the following basic syntax:

```
[ WITH ... ]  
SELECT  
    [ TOP <n> ]  
    ...  
[ FROM ...  
    [ AT | BEFORE ... ]  
    [ CHANGES ... ]  
    [ CONNECT BY ... ]  
    [ JOIN ... ]  
    [ MATCH_RECOGNIZE ... ]  
    [ PIVOT | UNPIVOT ... ]  
    [ VALUES ... ]  
    [ SAMPLE ... ] ]  
[ WHERE ... ]  
[ GROUP BY ...  
    [ HAVING ... ] ]  
[ QUALIFY ... ]  
[ ORDER BY ... ]  
[ LIMIT ... ]
```

## What is Snowflake SQL – SnowSQL?

**SnowSQL** is the Snowflake SQL command-line client that allows you to connect to Snowflake and execute SQL queries. It also lets you perform all DDL and DML operations, including loading and unloading of data from database tables.

SnowSQL (snowsql executable) can run as an interactive shell or in batch mode. This can be done by either using stdin or using the -f option.

**SnowSQL** was developed using the **Snowflake Connector for Python**. However, this connector is not required for installing SnowSQL. All the required software for installing SnowSQL is already bundled in the installers. SnowSQL is supported on all popular operating systems including Windows, macOS, and some distributions of Linux.

## Prerequisites:

1. Install and configure SnowSQL.
2. Your Snowflake user must have the necessary permissions to create different Snowflake objects.

## Logging into SnowSQL

1. Open a terminal window.
2. Start SnowSQL at the command prompt using the following command:

```
$ snowsql -a <accountName> -u <userName>
```

Here:

**<accountName>** is the name that has been assigned to your account by Snowflake.

**<userName>** is the login name assigned to your Snowflake user.

3. When you receive a prompt by SnowSQL, enter the password corresponding to your Snowflake user.

For more details, see **Connecting Through SnowSQL**.

## Creating Snowflake Objects

Additionally, loading and querying of data require a **virtual warehouse**. This provides the required compute resources to perform these tasks.

### 1. Create a Database

The database can be created using the following syntax:

```
CREATE OR REPLACE DATABASE demo;
```

Snowflake assigns a default schema named public to every database created and thus, it is not necessary to create a schema on your own. The database and schema that you just created are now being used for your current

session. This information is displayed in your SnowSQL command prompt, but you can also view it by using the following command:

```
SELECT CURRENT_DATABASE(), CURRENT_SCHEMA();
```

## 2. Creating a Table

Create a table named emp\_details in demo.public using the CREATE TABLE command:

```
CREATE OR REPLACE TABLE emp_details (  
  first_name STRING,  
  last_name STRING,  
  email STRING,  
  address STRING,  
  city STRING,  
  start_date DATE  
);
```

Note that the number of columns in the table, their positions, and their data types correspond to the fields in the CSV data files that you will be staging in the next step.

## 3. Creating a Virtual Warehouse

Create an X-Small warehouse named demo\_wh using the following syntax:

```
CREATE OR REPLACE WAREHOUSE demo_wh WITH  
  WAREHOUSE_SIZE='X-SMALL'  
  AUTO_SUSPEND = 180  
  AUTO_RESUME = TRUE  
  INITIALLY_SUSPENDED=TRUE;
```

Here,

**WAREHOUSE\_SIZE:** Specifies the number of servers in each cluster. To view the different sizes available, click [here](#).

**AUTO\_SUSPEND** and **AUTO\_RESUME:** It will suspend the warehouse if it is inactive for the specified period. Auto-resume has enabled automatically when a query is submitted that requires to compute resources from the warehouse and the warehouse is in the current session.

**INITIALLY\_SUSPENDED** = TRUE | FALSE, specifies whether the warehouse is created initially in the 'Suspended' state.

It should be noted that the warehouse will not be started initially, but is set to auto-resume, what it means is that, it will automatically start running when you execute your first SQL command that requires its compute resources.

To see other parameters and allowed values for creating a warehouse, click [here](#).

Also, note that the warehouse you created is now being used for your current session. This information will be displayed in your SnowSQL command prompt but can also be viewed using the following command:

```
SELECT CURRENT_WAREHOUSE();
```

## Staging Data Files

Snowflake lets you load data from files that have been not only been staged in an internal (Snowflake) stage but also in an external (Amazon S3, Google Cloud Storage, or Microsoft Azure) stage. Loading from any external stage is very useful when you already have data files stored in these cloud storage services.

The **PUT** command is used to stage files.

Execute **PUT to upload local data files to the table stage** provided for the emp\_details table you previously created. Note that this command is OS-specific because it references files in your local environment. The path thus needs to be set accordingly.

### Linux or macOS

```
PUT file:///tmp/employees0*.csv @demo.public.%emp_details;
```

### Windows

```
PUT file:///C:/tempemployees0*.csv @demo.public.%emp_details;
```

It is important to note that the **PUT** command **compresses files** by default using **gzip**.

You can see the list of the files you successfully staged by executing a **LIST** command:

```
LIST @demo.public.%emp_details;
```

## Copy Data into the Target Table

You can now easily load your staged data into the target table using the **COPY INTO <table>** command. Another important thing to keep in mind is that this command requires an active and running warehouse.

```
COPY INTO emp_details
FROM @%emp_details
FILE_FORMAT = (type = csv field_optionally_enclosed_by='')
PATTERN = '*.employees0[1-5].csv.gz'
ON_ERROR = 'skip_file';
```

Let's look more closely at this command:

- The **FROM** clause is used to identify the internal stage location.
- **FILE\_FORMAT** is used to specify the type of the file. In this example, it specifies the file type as CSV and the double-quote character (") as the character used to enclose strings. Snowflake supports a variety of file types and options.
- **PATTERN** is used for **pattern matching to pull data** from all files that match the regular expression `*employees0[1-5].csv.gz`.
- **ON\_ERROR** specifies what to do when the COPY command incurs any errors in the files. By default, the command stops loading data when the first error is caught. However, we have edited the query to skip any file containing an error and move on to loading the next file.

Moreover, The **COPY** command helps in **validating files before loading**. You can go through the **COPY INTO <table>** topic for additional error checking and validation methods.

## Querying the Loaded Data

The data loaded in the emp\_details table can be queried easily using the standard Snowflake **SQL** statements, supporting **functions**, and **operators**.

Moreover, the data can also be manipulated, like updating the loaded data or inserting more data, using standard **DML commands**.

1. To return all rows and columns from the table:

```
SELECT * FROM emp_details;
```

2. To insert rows directly into a table:

E.g. to insert two additional rows into the table:

```
INSERT INTO emp_details VALUES  
  ('Clementine', 'Adamou', 'cadamou@sf_tuts.com', '10510 Sachs  
Road', 'Klenak', '2017-9-22') ,  
  ('Marlowe', 'De Anesy', 'madamouc@sf_tuts.co.uk', '36768  
Northfield Plaza', 'Fangshan', '2017-1-26');
```

3. For pattern matching, you can use the LIKE function.

E.g. To display records having email addresses with a UK domain name:

```
SELECT email FROM emp_details WHERE email LIKE '%.uk';
```